


Project5

January 23, 2021

1 Brief Review

The structure of the layers of VGG-16 Model which is trained on the ImageNet Dataset is shown below.

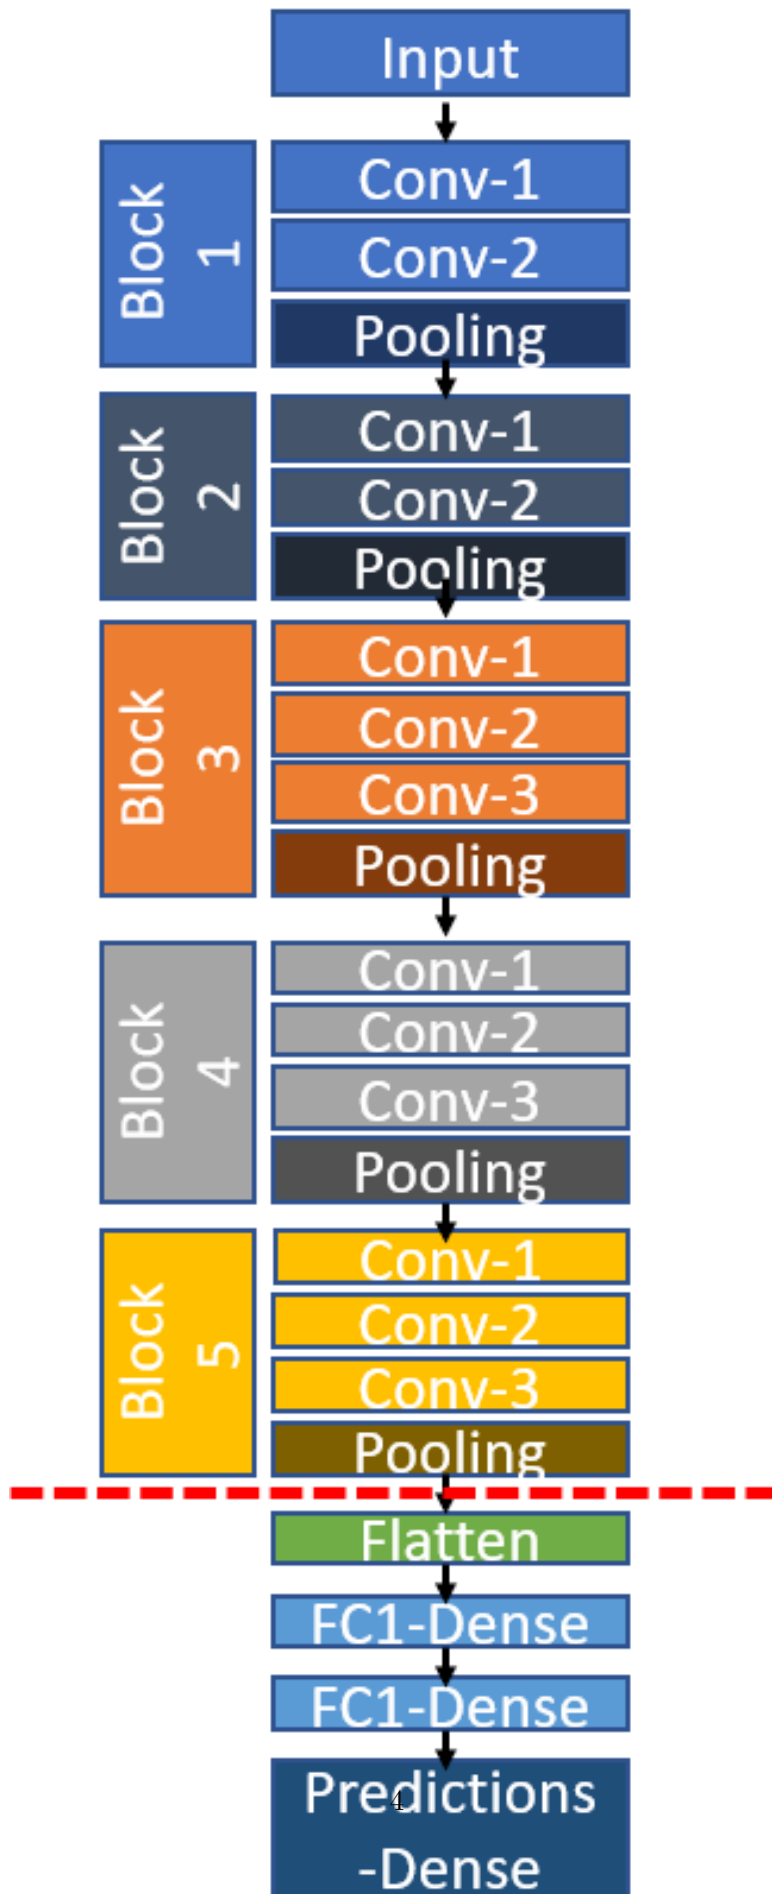


And also you can see another table which describes the summary of VGG-16

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
flatten (Flatten)	(None, 25088)	0
fc1 (Dense)	(None, 4096)	102764544
fc2 (Dense)	(None, 4096)	16781312
predictions (Dense)	(None, 1000)	4097000

The VGG-16 contains 16 layers. As you can see, after flatten of the last convolutional layer, there is 25,088 features (the 1st highlighted) and in the final layer (prediction or final dense layer), there are 1000 nodes (the 2nd highlighted), because VGG-16 is trained for 1000-class classification problems.

For using this Model and applying transfer learning on that for our specific 6-class classification, we can extract 25088 features from the Medical-MNIST dataset using the VGG model. The cut for feature extraction has been shown in the image below.



1.1 VGG-16 model

```
[2]: from keras.layers import Input, Lambda, Dense, Flatten
from keras.models import Model
from keras.applications.vgg16 import VGG16
from keras.utils import plot_model

model = VGG16()
model.summary()
```

Model: "vgg16"

Layer (type)	Output Shape	Param #
input_2 (InputLayer)	(None, 224, 224, 3)	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808

block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808

block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0

flatten (Flatten)	(None, 25088)	0

fc1 (Dense)	(None, 4096)	102764544

fc2 (Dense)	(None, 4096)	16781312

predictions (Dense)	(None, 1000)	4097000
=====		
Total params: 138,357,544		
Trainable params: 138,357,544		
Non-trainable params: 0		

```
[4]: from keras import models

base_model = VGG16(weights='imagenet')
model_VGG16 = models.Model(inputs=base_model.input, outputs=base_model.
    ↳get_layer('flatten').output)
model_VGG16.summary()
```

Model: "model_2"

Layer (type)	Output Shape	Param #
=====		
input_4 (InputLayer)	(None, 224, 224, 3)	0

block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792

block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928

block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0

block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856

block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584

block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0

block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168

block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080

block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080

block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0

block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160

block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808

block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808

block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0

block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808

block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808

block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808

block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0

flatten (Flatten)	(None, 25088)	0
=====		
Total params: 14,714,688		
Trainable params: 14,714,688		
Non-trainable params: 0		

2 Importing Medical-MNIST dataset

Now we will import the images from Medical-MNIST dataset and convert them to a normalized value for each pixel and a number between 0 to 5 for each class.

```
[5]: import numpy as np
import pandas as pd
import os
import torch
from PIL import Image

torch.manual_seed(1)
classes=['Abdomen CT',"Breast MRI",'Chest XRay','Chest CT','Hand','Head CT']

AbdomenCT_dir='./medical-mnist/AbdomenCT/'
BreastMRI_dir='./medical-mnist/BreastMRI/'
ChestXRay_dir='./medical-mnist/CXR/'
ChestCT_dir='./medical-mnist/ChestCT/'
Hand_dir='./medical-mnist/Hand/'
HeadCT_dir='./medical-mnist/HeadCT/'
```

```

directories=[AbdomenCT_dir, BreastMRI_dir, ChestXRay_dir, ChestCT_dir,
↪Hand_dir, HeadCT_dir]

def create_dataset(directories, height=64,width=64):
    X=[]
    Y=[]
    data_index = 0
    for class_num in range(len(directories)):
        print(class_num)
        images=os.listdir(directories[class_num])
        print(images[0])
        print(len(images))
        print(directories[class_num]+images[0])
        for i in range(len(images)):
            images[i] = directories[class_num]+images[i]
            image_x = Image.open(images[i])
            image_x = image_x.resize((width,height))
#             image_x = image_x.convert('L')
            image_x = np.asarray(image_x)
            image_x = image_x.astype('float32')
            image_x /= 255
            image_y = class_num
            X.append(image_x)
            Y.append(image_y)
    return X, Y

X, y = create_dataset(directories=directories)

```

```

0
001498.jpeg
10000
./medical-mnist/AbdomenCT/001498.jpeg
1
001498.jpeg
8954
./medical-mnist/BreastMRI/001498.jpeg
2
001498.jpeg
10000
./medical-mnist/CXR/001498.jpeg
3
001498.jpeg
10000
./medical-mnist/ChestCT/001498.jpeg
4
001498.jpeg

```



```

10000
./medical-mnist/Hand/001498.jpeg
5
001498.jpeg
10000
./medical-mnist/HeadCT/001498.jpeg

```

3 Adding our preferred flatten layer to the VGG model

```

[7]: from keras import optimizers
     for layer in model.layers:
         layer.trainable = False

     prediction = Dense(6, activation='softmax')(model_VGG16.output)
     model = Model(inputs=model_VGG16.input, outputs=prediction)
     model.summary()

```

Model: "model_4"

Layer (type)	Output Shape	Param #
input_4 (InputLayer)	(None, 224, 224, 3)	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808

block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808

block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0

block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808

block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808

block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808

block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0

flatten (Flatten)	(None, 25088)	0

dense_2 (Dense)	(None, 6)	150534
=====		
Total params: 14,865,222		
Trainable params: 150,534		
Non-trainable params: 14,714,688		

```
[ ]: from keras import optimizers

model.compile(loss='categorical_crossentropy',
              optimizer=optimizers.Adam(),
              metrics=['accuracy'])

X = np.array((X))
y = np.array((y))

model.fit(X, y, batch_size=32, epochs=20, validation_split=0.3)
```